

Operating Systems

(Sistemas Operativos)

CPU Virtualization Mechanisms

University of Minho
2024-2025



What will we learn?

CPU Management

- ◎ Although a computer has a limited number of CPUs, users get the impression that several programs (a lot more than the available CPUs) are running simultaneously
 - How does the OS provide this illusion?
 - How can users use the OS APIs to start and stop programs?
 - How does the OS choose what programs will be running, be switched, ...?
- ◎ Let us start with the **mechanisms!**

CPU Virtualization

Challenges

- To virtualize the CPU, the OS must share it among many programs running seemingly at the same time
 - Basic idea: **Time-sharing** the CPU
i.e., run a process for a little while, then another one, ...
- Challenges
 - **Performance:** OS mechanisms cannot add excessive overhead to the system
 - **Control:** The OS must retain control over the CPU, avoiding processes from
 - Taking control of the CPU (e.g., infinite loops)
 - Accessing/corrupting information from other programs in memory or at disk

Limited Direct Execution

Basic Technique


● **Goal:** Run a program as fast as possible in the CPU (Performance)

- **Solution:** The program runs natively on the hardware CPU (*direct execution*)

● **Issues:** If the program has full control over the CPU, how does the OS...

- Ensures the process does not perform any unwanted operations?
- Stops the process from running and switches it with another one?

● One needs to *limit the execution* of programs (*i.e.*, limit what these can do)



Modes of Execution

Restricted Operations


- User processes (*i.e.*, running programs) cannot freely do restricted operations (*e.g.*, I/O requests to disk, gain access to resources such as CPU or memory)
- The OS must mediate this interaction (with the aid of hardware)!
The CPU provides two modes of execution¹
 - **User mode:** Code running in the CPU cannot do restricted (privileged) operations. (*e.g.*, a process in user mode cannot directly do disk I/O operations)
 - **Kernel mode:** Code executing in the CPU has full access to the hardware and can do any restricted operation. **The OS runs in this mode!**
- **Issues:** What if a user process wishes to do a privileged operation?
 - Remember the system calls?

¹A single bit, stored at the processor, indicates which mode the CPU is currently running in

System Calls

Restricted Operations

- **System calls** allow the kernel (OS) to expose privileged operations to user processes (e.g., process creation, communication, memory allocation, I/O)
- To execute a system call, the process calls a special **trap** instruction, raising the privilege level to **kernel mode** and jumping into the OS code (**trap handler**)
- The OS can now perform privileged operations (on behalf of the user process)
- When finished the OS calls a **return-from-trap** instruction that reduces the privilege level to **user mode** and resumes the execution of the process



Limited Direct Execution (Updated)

OS (Kernel Mode)

Hardware

Process (User Mode)

Limited Direct Execution (Updated)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

Hardware

Process (User Mode)

Limited Direct Execution (Updated)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

Hardware

move to user mode
jump to main

Process (User Mode)

Limited Direct Execution (Updated)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

Hardware

move to user mode
jump to main

Process (User Mode)

run *main()* ...
call system call
trap into OS

Limited Direct Execution (Updated)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

Hardware

move to user mode
jump to main

move to kernel mode
jump to trap handler

Process (User Mode)

run *main()* ...
call system call
trap into OS

Limited Direct Execution (Updated)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

handle trap
do work of system call
return-from-trap

Hardware

move to user mode
jump to main

move to kernel mode
jump to trap handler

Process (User Mode)

run *main()* ...
call system call
trap into OS

Limited Direct Execution (Updated)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

handle trap
do work of system call
return-from-trap

Hardware

move to user mode
jump to main

move to kernel mode
jump to trap handler

move to user mode
jump to PC after trap

Process (User Mode)

run *main()* ...
call system call
trap into OS

Limited Direct Execution (Updated)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

handle trap
do work of system call
return-from-trap

Hardware

move to user mode
jump to main

move to kernel mode
jump to trap handler

move to user mode
jump to PC after trap

Process (User Mode)

run *main()* ...
call system call
trap into OS

...
return from main
trap (via *exit()*)

Limited Direct Execution (Updated)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

handle trap
do work of system call
return-from-trap

free process memory / remove PCB

Hardware

move to user mode
jump to main

move to kernel mode
jump to trap handler

move to user mode
jump to PC after trap

Process (User Mode)

run *main()* ...
call system call
trap into OS

...
return from main
trap (via *exit()*)

Mode Switching

Traps and Interrupts

- ◎ **Two types of instructions** lead to transferring control from user into kernel mode
- ◎ A **trap** is generated by the current process running on a processor
 - The user program asks for an OS service (as we saw before)
e.g., when the process issues a system call (in UNIX)
 - The user program performs an illegal action (an exception)
e.g., division by zero, illegal memory access, ...
- ◎ **Interrupt** are caused by devices and may not be related with the currently running process
 - When a device (disk, keyboard, network card) requires the OS attention
e.g., a disk read is ready, there is keyboard input to be read

Note: Terminology changes depending on the system (x86, xv6). Some call trap to all instructions, other use the terms interchangeably...

Interrupts

Why are interrupts useful?

- ◎ **Example:** When a process issues a read for a file (system call), the OS will ask the disk device for that specific content
 - The disk will require some time to bring that information to memory...
- ◎ The OS can periodically check (**polling**) the disk device for that request completion
 - If events take long time to complete, then the OS is waisting precious CPU cycles that could be used to run other processes!
- ◎ **Interrupts** are generated by the devices asynchronously
 - **Benefit:** The OS is notified when events occur, leaving the CPU free for other processes
 - **Caution:** Interrupts are issued concurrently, when other processes or the OS are running at the CPU

Limited Direct Execution (Interrupts)

OS (Kernel Mode)

| create PCB entry / allocate and load memory
| **return-from-trap**

Hardware

| move to user mode
| jump to main

Process (User Mode)

| run *main()*
| ...

Limited Direct Execution (Interrupts)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

Hardware

move to user mode
jump to main

interrupt event (e.g., from disk)
move to kernel mode
jump to trap handler

Process (User Mode)

run *main()*
...

Limited Direct Execution (Interrupts)

OS (Kernel Mode)

create PCB entry / allocate and load memory
return-from-trap

handle interruption
do work
return-from-interrupt

free process memory / remove PCB

Hardware

move to user mode
jump to main

interrupt event (e.g., from disk)
move to kernel mode
jump to trap handler

move to user mode
jump to PC after interruption

Process (User Mode)

run *main()*
...

...
return from main
trap (via *exit()*)

Traps and Interrupts

Process state

- Both **trap** and **interrupt** instructions are **hardware-assisted**
 - The OS cannot take the CPU from a process without aid
 - The CPU must interrupt the process, change to kernel mode, save some information and jump to the OS code
- When a process is interrupted, the CPU must save a set of registers
 - Registers contain information (e.g., program counter) required by the CPU to resume the process execution when the OS calls a **return-from-trap/interrupt** instruction
 - Registers are stored, for example, in a per-process kernel stack

*On x86, for example, the processor will push the **program counter, flags, and a few other registers** onto its **kernel stack**; the return-from-trap will pop these values off the stack and resume execution of the user-mode program*

Limited Direct Execution (Registers)

OS (Kernel Mode)

Hardware

Process (User Mode)

Limited Direct Execution (Registers)

OS (Kernel Mode)

create PCB / allocate and load memory
fill kernel stack with registers / PC
return-from-trap

Hardware

Process (User Mode)

Limited Direct Execution (Registers)

OS (Kernel Mode)

create PCB / allocate and load memory
fill kernel stack with registers / PC
return-from-trap

Hardware

restore regs (from kernel stack)
move to user mode
jump to main

Process (User Mode)

Limited Direct Execution (Registers)

OS (Kernel Mode)

create PCB / allocate and load memory
fill kernel stack with registers / PC
return-from-trap

Hardware

restore regs (from kernel stack)
move to user mode
jump to main

Process (User Mode)

run *main()* ...
call system call / **trap** into OS

Limited Direct Execution (Registers)

OS (Kernel Mode)

- create PCB / allocate and load memory
- fill kernel stack with registers / PC**
- return-from-trap**

Hardware

- restore regs (from kernel stack)**
- move to user mode
- jump to main

- save regs (to kernel stack)**
- move to kernel mode
- jump to trap handler

Process (User Mode)

- run *main()* ...
- call system call / **trap** into OS

Limited Direct Execution (Registers)

OS (Kernel Mode)

create PCB / allocate and load memory
fill kernel stack with registers / PC
return-from-trap

handle trap
do work of system call
return-from-trap

Hardware

restore regs (from kernel stack)
move to user mode
jump to main

save regs (to kernel stack)
move to kernel mode
jump to trap handler

Process (User Mode)

run *main()* ...
call system call / **trap** into OS

Limited Direct Execution (Registers)

OS (Kernel Mode)

create PCB / allocate and load memory
fill kernel stack with registers / PC
return-from-trap

handle trap
do work of system call
return-from-trap

Hardware

restore regs (from kernel stack)
move to user mode
jump to main

save regs (to kernel stack)
move to kernel mode
jump to trap handler

restore regs (from kernel stack)
move to user mode
jump to PC after trap

Process (User Mode)

run *main()* ...
call system call / **trap** into OS

Limited Direct Execution (Registers)

OS (Kernel Mode)

create PCB / allocate and load memory
fill kernel stack with registers / PC
return-from-trap

handle trap
do work of system call
return-from-trap

Hardware

restore regs (from kernel stack)
move to user mode
jump to main

save regs (to kernel stack)
move to kernel mode
jump to trap handler

restore regs (from kernel stack)
move to user mode
jump to PC after trap

Process (User Mode)

run *main()* ...
call system call / **trap** into OS

...
return from main / **trap** (via *exit()*)

Limited Direct Execution (Registers)

OS (Kernel Mode)

create PCB / allocate and load memory
fill kernel stack with registers / PC
return-from-trap

handle trap
do work of system call
return-from-trap

Free process memory / remove PCB

Hardware

restore regs (from kernel stack)
move to user mode
jump to main

save regs (to kernel stack)
move to kernel mode
jump to trap handler

restore regs (from kernel stack)
move to user mode
jump to PC after trap

Process (User Mode)

run *main()* ...
call system call / **trap** into OS

...
return from main / **trap** (via *exit()*)

Limited Direct Execution (Registers)

OS (Kernel Mode)

create PCB / allocate and load memory
fill kernel stack with registers / PC
return-from-trap

Hardware

restore regs (from kernel stack)
move to user mode
jump to main

Process (User Mode)

run *main()* ...
call system call / **trap** into OS

save regs (to kernel stack)
move to kernel mode
jump to trap handler

handle trap
do work of system call
return-from-trap



How does the hardware knows what OS code it should run (it should jump into)?

Traps and Interrupts

Interrupt/Trap Table

- At boot time the kernel (OS) sets up an **interrupt/trap table**
- The table tells the hardware what code to run (**interrupt/trap handler**) under exceptional events
 - what code should run when an hard disk interrupt happens,
 - or when a keyboard interrupt occurs,
 - or when a program makes a system call
- The table maps interrupts and traps (e.g., identified by a number) to their corresponding handlers
 - It distinguishes interrupts from different **devices** (e.g., disk, keyboard) and **requests** (e.g, disk read, write)
 - It distinguishes between different **trap types** (e.g., exception, system call)

0	HANDLER SYSTEM CALL
1	HANDLER KEYBOARD...
2	HANDLER DISK...
...	...

Illustrative example
(not a real one!)

System Calls in Practice

Integration with the C library

- When a program calls a system call (e.g., *fork()*, *open()*), a procedure call into the C library is done
- The library uses an agreed-upon calling convention with the OS kernel
 - **System call arguments** and **number** (identifier) are placed into well-known locations (e.g., stack, special registers)
- The library executes the corresponding trap instruction (system call)
 - The **interrupt/trap table** is used to check for the **system call trap handler**
 - When running the trap handler code, the OS can now validate the system call **number** and run the corresponding code
- This indirection provides **security**
(safer than jumping directly to a memory address to run the code)

Switching Processes

The Conundrum

- ◎ Let us go back to our initial goal
 - **Time-sharing** the CPU, i.e., run a process for a little while, then another one...
 - This requires pausing one process and running another one
 - But we saw that if a process is running at the CPU, the OS is not running
- ◎ **Question:** Given the mechanisms studied before, do you have any ideas on how the OS can share the CPU across several processes?
- ◎ We are getting closer to revealing the illusion created by the OS!

Cooperative Approach

Wait for System Calls

- When a process does a system call, the OS kernel regains control to
 - Handle the system call (e.g., ask the disk to read some content)
 - Do other tasks such as scheduling another process!
- Actually the OS can also regain control under device interruptions and other traps (exceptions)
- In a cooperative approach the OS assumes the processes will
 - Issue system calls (e.g., block for doing disk I/O)
 - or voluntarily yield the CPU to the OS (the *yield* system call serves exactly this purpose)
- **Question:** What if we have malicious or buggy programs (e.g., infinite loops)...?
Any ideas on how to solve it?

Non-Cooperative Approach

Timer Interrupt

- ◎ We need help from the hardware once again, namely from a **timer interrupt**
 - *E.g.*, every X milliseconds an interrupt is generated by the hardware
 - The process is halted, a specific interrupt handler runs and the OS gains control!
- ◎ The OS is responsible for
 - Defining the code for the timer interrupt handler
 - Setting the timer's interval (at boot time)
- ◎ **Remember:** when the interrupt occurs process state must be saved so the process can be resumed later...

Saving and Restoring Context

Context Switching

- ◎ When the OS regains control (cooperatively or non-cooperatively)
 - it must decide to keep running the same process or to run another one
 - This decision is made by the OS **scheduler**
- ◎ When the scheduler decides to switch process_A with process_B, a **context switch** is made
 - The OS must save information of process_A (e.g., program counter, memory limits, kernel stack pointer¹) into its PCB_A structure (**Process Control Block** - check the *Processes slides*!)
 - This information is part of the **process list** and it is managed by the OS!
(Different from the per-process kernel stack managed by the hardware)
 - Switches contexts, by changing the CPU's stack pointer (a register) to the kernel stack of process_B

¹ Remember that the CPU stores essential information to resume a process in a per-process kernel stack. The PCB stores a pointer to the stack's address.

Limited Direct Execution (Context Switching)

OS (Kernel Mode)

Hardware

Processes (User Mode)

Limited Direct Execution (Context Switching)

OS (Kernel Mode)

Hardware

Processes (User Mode)

Process A

...

Limited Direct Execution (Context Switching)

OS (Kernel Mode)

Hardware

Processes (User Mode)

Process A

...

timer interrupt

save regs (to kernel stack_A)

move to kernel mode

jump to interrupt handler

Limited Direct Execution (Context Switching)

OS (Kernel Mode)

Hardware

Processes (User Mode)

Process A

...

timer interrupt
save regs (to kernel stack_A)
move to kernel mode
jump to interrupt handler

handle interrupt
call *switch()* routine
save PCB_A
restore regs from PCB_B
switch to kernel stack_B

Limited Direct Execution (Context Switching)

OS (Kernel Mode)

Hardware

Processes (User Mode)

Process A

...

timer interrupt

save regs (to kernel stack_A)

move to kernel mode

jump to interrupt handler

handle interrupt

call **switch()** routine

save PCB_A

restore regs from PCB_B

switch to kernel stack_B

restore regs (from kernel stack_B)

move to user mode

jump to Process_B PC

Limited Direct Execution (Context Switching)

OS (Kernel Mode)

Hardware

Processes (User Mode)

Process A

...

timer interrupt
save regs (to kernel stack_A)
move to kernel mode
jump to interrupt handler

handle interrupt
call *switch()* routine
save PCB_A
restore regs from PCB_B
switch to kernel stack_B

restore regs (from kernel stack_B)
move to user mode
jump to Process_B PC

Process B

...

Performance

Context Switching

- ◎ In 1996, when running Linux 1.3.37 on a 200-MHz P6 CPU, **system calls took roughly 4 microseconds**, and a **context switch roughly 6 microseconds**
- ◎ Modern systems perform almost an order of magnitude better, with **sub-microsecond** results on systems with 2- or 3-GHz processors
- ◎ Whenever a system call to disk is made, for example, the OS will probably schedule another process to run
 - **Question:** Why?
 - **Question:** Do you see an impact of implementing a program that does repetitive system calls that take few microseconds?

More Information

- **Chapters 6** - Remzi H. Arpaci-Dusseau, Andrea C. Arpaci-Dusseau. **Operating Systems: Three Easy Pieces**. Arpaci-Dusseau Books, 2018.
- Avi Silberschatz, Peter Baer Galvin, Greg Gagne. **Operating System Concepts (10. ed)**. John Wiley & Sons, 2018.

Questions?